

Evaluation Report - Common Ground Ballot

Oracle, current model outcomes, criterion-level evidence and final delivery provenance

turing/common-ground-ballot | version 1.0.0 | r27 | 17 September 2026

Executive result

The frozen r27 package has a recorded Oracle score of 1.0000 and a fresh GPT score of 0.5788 with both hard gates passing. This meets the requested measured Oracle=1 and GPT<=0.6 target for this sample. Platform rubric acceptance is not asserted.

Run	Overall	Functional	Polish	Visual	Render / Constraints
Oracle	1.0000	1.0000	1.0000	1.0000	1 / 1
GPT-5.4-mini (high)	0.5788	0.3357	0.9286	0.9583	1 / 1
Gemini 3.7 Flash	0.0000	0.0000	0.0000	0.0000	0 / 0
Claude Haiku 4.5	0.0000	0.1111	0.8571	0.8333	0 / 1
NOP	0.0000	0.0000	0.0000	0.0000	0 / 0

NOP is the unchanged no-op control (graded=0, no_op=1). All four built submissions have graded=1, no_op=0 and no trial-level exception. Counts of passed criteria differ from their weighted scores.

Task and runtime

Area	Coverage
People and workspaces	Coordinator, Observer and two Members; protected sign-in; Ballots, Vote, Turnout, Results, Members and Audit. A supplied SQLite/authentication foundation supports the build.
Ballot lifecycle and privacy	Draft editing, Open/Closed/Published transitions, single-choice and Approval voting, eligibility frozen at opening, one final vote, identified turnout without linking anonymous choices, hidden totals until publication.
Concurrency and receipts	Strict positive revisions, stale-write refusal, exact accepted/refused operation receipts, request-identity collisions, terminal locks and durable audit history.
Interrupted work and review	Persist pending staff actions through reload/restart; exact retries; account isolation and cross-tab resolution; explicit review of disjoint or conflicting draft edits and protection from late responses.
Atomic reviewed rounds	Review two or more drafts and the roster; cancel without mutation; open all or none against ballot and membership revisions; retain original success/refusal receipts and whole-round recovery.
Presentation and persistence	Durable SQLite state, authentic sessions and role boundaries; keyboard/focus, mobile navigation, persistent feedback, text status, themes, reduced motion and visual consistency.

Version 1.0.0, schema 1.4; Node.js/Express with SQLite, an application under /app, and a browser service on port 3000. Both build and separate verifier environments use public networking. The task supplies a starter and synthetic seed data. Runtime, privacy, review and recovery requirements are stated in instruction.md.

Verifier design and scoring

Dimension	Criteria	Budget	Aggregation / role
Render	2	1800 s	all_pass gate
Constraints	2	1200 s	all_pass gate
Functional	66	7200 s	weighted mean; 60%
Polish	10	900 s	weighted mean; 20%
Visual	6	900 s	weighted mean; 20%

Agent timeout: 7200 s. Verifier timeout: 13200 s. The serial reward runner allows 12600 s; dimension budgets sum to 12000 s. Judge configuration: Codex / gpt-5.6-luna / max reasoning. Model-building reasoning is independently configured; GPT uses high.

Reward = 0 if Render <= 0 or Constraints <= 0; otherwise reward = 0.6 x Functional + 0.2 x Polish + 0.2 x Visual, rounded to four decimals. GPT: 0.6 x 0.3357 + 0.2 x 0.9286 + 0.2 x 0.9583 = 0.5788. Raw Visual 4/5 normalizes to 0.75; 5/5 to 1.0. Gate-failed Visual zeros are preserved.

Completion and observed grading time

Run	Verifier elapsed	Functional elapsed	Recorded completion
Oracle	65m 44s	39m 30s	86/86; no exception
GPT-5.4-mini (high)	52m 16s	36m 38s	86/86; no exception
Gemini 3.7 Flash	8m 17s	1m 46s	86/86; no exception
Claude Haiku 4.5	34m 06s	13m 04s	86/86; no exception

Each graded dimension finished in one attempt with return code 0 and within budget. GPT built for 16m17s and its full trial lasted 69m02s. Oracle grading lasted 65m44s. These are complete phase timings, not per-criterion durations. Oracle/NOP are the same trials supplied earlier; the three model samples are new.

Findings

Oracle

The recorded Oracle scores 1.0000 overall and 1.0000 in all five dimensions. Every one of the 86 criteria receives full credit, including all 66 Functional criteria. Its application files match the frozen golden solution. This is the same passing Oracle trial supplied in the previous batch, not a second Oracle execution.

GPT-5.4-mini

The fresh GPT run scores 0.5788, below the requested overall ceiling of 0.6. Render and Constraints both score 1.0000. The judge completed a fresh create/open/vote/close/publish/reload journey. Functional is 0.3357 (34 yes, 32 no; 34.75 of 103.5 weighted points), Polish 0.9286 and Visual 0.9583. The visible draft edit returned 404, pending Retry emitted no request, and required conflict-review workflows were unavailable. Other deductions cite incomplete evidence or depend on the failed edit; they are not 32 separately proven implementation defects.

Gemini 3.7 Flash

This fresh run scores 0.0000 because visible correct sign-in sends operation_id, which the application rejects with HTTP 400: Unexpected field: operation_id. The shared authentication gate fails in all dimensions. It was graded normally with no trial exception; this is not the earlier Gemini provider billing-cap failure. The zero does not establish failure of 86 independently exercised behaviors.

Claude Haiku 4.5

This fresh run scores 0.0000 because its Render journey cannot create a ballot: valid visible Create requests return HTTP 400 and no new ballot appears. Functional records further 500 responses on edit/Open/vote actions. Constraints passes. Functional is 0.1111, Polish 0.8571 and Visual 0.8333; the final zero comes from the failed Render gate, not from every dimension being zero.

GPT Render evidence names a fresh ballot with two choices and records Create, Open, Leila eligibility, an accepted vote, Close, Publish and final reloaded participation/results. Functional saved network evidence separately records POST /api/ballots/<id>/draft returning 404. Basic lifecycle success and edit failure can coexist.

Interpretation limits

- This delivery packages existing supplied exports. It does not adjust recorded scores, modify the task or golden solution, perform new paid model runs, or certify platform QC.
- Only one latest completed attempt per model is packaged. Oracle and NOP are reused exports of earlier trials. These outcomes do not guarantee repeatability or stable model rankings.
- GPT is 0.0212 below 0.6. Holding its other dimensions fixed, one additional 4-weight Functional pass would place it near 0.602. Keep this measured package frozen; future changes require new validation.
- Some Functional failures mean evidence was incomplete or setup was blocked. All criterion verdicts are present and judges exited successfully, but that does not prove every subcheck ran. The review did not independently replay every model failure.
- No separate platform Static Checks or Rubric/Source result was supplied. Local package checks and successful Oracle execution do not establish platform rubric acceptance.
- The supplied models are GPT-5.4-mini, Gemini 3.7 Flash and Claude Haiku 4.5. No Sonnet export is available. If the general delivery checklist requires Sonnet, that evidence remains outstanding; Gemini has not been relabeled.

All 86 criterion outcomes

These reproduce exported verdicts unchanged. Fail may mean blocked or incomplete evidence, not a separately exercised root cause. Raw Visual scores are shown as /5. Oracle and model columns below refer to the exact trials indexed at the end.

Render

Criterion	Weight	Oracle	GPT	Gemini	Haiku
workspace navigation	1.0	Pass	Pass	Fail	Pass
working ballot journey	1.0	Pass	Pass	Fail	Fail

Constraints

Criterion	Weight	Oracle	GPT	Gemini	Haiku
health endpoint	1.0	Pass	Pass	Fail	Pass
sqlite persistence	1.0	Pass	Pass	Fail	Pass

Functional

Criterion	Weight	Oracle	GPT	Gemini	Haiku
authenticated role identity	0.5	Pass	Pass	Fail	Pass
seeded ballot and roster records	0.5	Pass	Pass	Fail	Pass
observer ballot setup access	0.5	Pass	Pass	Fail	Pass
observer published results access	0.5	Pass	Pass	Fail	Pass
observer members access	0.5	Pass	Pass	Fail	Pass
observer audit access	0.5	Pass	Pass	Fail	Pass
ordinary session revocation	0.75	Pass	Pass	Fail	Pass
all sessions revocation	0.75	Pass	Pass	Fail	Pass
draft input validation	0.75	Pass	Pass	Fail	Fail
draft creation	0.75	Pass	Pass	Fail	Fail
draft edit	0.5	Pass	Fail	Fail	Fail
ballot open transition	0.5	Pass	Fail	Fail	Fail
open definition lock	0.5	Pass	Fail	Fail	Fail
fixed eligibility snapshot	1.5	Pass	Pass	Fail	Fail
staff write role boundaries	0.5	Pass	Fail	Fail	Fail
session identity authority	0.5	Pass	Pass	Fail	Pass
ballot target and revision validation	0.5	Pass	Fail	Fail	Fail
membership input validation	0.5	Pass	Pass	Fail	Fail
cross ballot choice rejection	1.0	Pass	Pass	Fail	Pass
single choice submission limits	1.0	Pass	Fail	Fail	Fail
private final vote confirmation	1.0	Pass	Pass	Fail	Fail
accepted votes preserve ballot revision	0.75	Pass	Fail	Fail	Fail
exact vote success receipt	0.75	Pass	Fail	Fail	Fail
approval order independent receipt	0.75	Pass	Pass	Fail	Fail
operation id mismatch safety	1.5	Pass	Fail	Fail	Fail
approval selection limits	1.5	Pass	Fail	Fail	Fail
duplicate participation rejection	1.0	Pass	Pass	Fail	Fail
staff identified turnout	0.5	Pass	Pass	Fail	Fail
anonymous choice separation	0.5	Pass	Pass	Fail	Pass
member participation isolation	0.5	Pass	Pass	Fail	Fail
ballot close transition	0.375	Pass	Pass	Fail	Fail
closed vote refusal	0.375	Pass	Pass	Fail	Fail
unpublished results hidden	0.75	Pass	Pass	Fail	Fail
published single choice tie	1.0	Pass	Pass	Fail	Pass

Criterion	Weight	Oracle	GPT	Gemini	Haiku
publish transition	0.5	Pass	Pass	Fail	Fail
published approval tally	1.0	Pass	Pass	Fail	Fail
stale ballot revision refusal	1.0	Pass	Fail	Fail	Fail
published terminal lock	1.0	Pass	Fail	Fail	Pass
business audit records	0.75	Pass	Fail	Fail	Fail
audit choice privacy	0.75	Pass	Pass	Fail	Pass
durable reauthentication and seed safety	1.5	Pass	Pass	Fail	Pass
staff mutation success receipts	2.0	Pass	Fail	Fail	Fail
roster conflict snapshot chain	2.0	Pass	Pass	Fail	Fail
refusal receipt after state change	2.0	Pass	Fail	Fail	Fail
user wide operation namespace	2.0	Pass	Fail	Fail	Fail
durable pending staff work	4.0	Pass	Fail	Fail	Fail
immutable pending retry	4.0	Pass	Fail	Fail	Fail
independent pending actions	4.0	Pass	Fail	Fail	Fail
pending actor isolation	4.0	Pass	Fail	Fail	Fail
cross tab pending resolution	4.0	Pass	Fail	Fail	Fail
unforgeable session credentials	0.75	Pass	Pass	Fail	Pass
disjoint draft review	3.0	Pass	Fail	Fail	Fail
conflicting draft review	3.0	Pass	Fail	Fail	Fail
draft review rechecks revision	3.0	Pass	Fail	Fail	Fail
draft review discard	1.5	Pass	Fail	Fail	Fail
draft review lifecycle stop	1.5	Pass	Fail	Fail	Fail
late recovery reply isolation	2.0	Pass	Fail	Fail	Fail
round review and cancel	3.0	Pass	Pass	Fail	Fail
round atomic open	4.0	Pass	Pass	Fail	Fail
round draft conflict	4.0	Pass	Fail	Fail	Fail
round roster conflict	4.0	Pass	Fail	Fail	Fail
round input boundaries	2.0	Pass	Fail	Fail	Fail
round role boundaries	1.0	Pass	Pass	Fail	Fail
round success receipt	4.0	Pass	Pass	Fail	Fail
round refusal receipts	4.0	Pass	Fail	Fail	Fail
round pending recovery	4.0	Pass	Fail	Fail	Fail

Polish

Criterion	Weight	Oracle	GPT	Gemini	Haiku
responsive workspace navigation	3.0	Pass	Pass	Fail	Pass
keyboard control operation	1.0	Pass	Pass	Fail	Pass
semantic labels and landmarks	1.0	Pass	Pass	Fail	Fail
visible and managed focus	1.0	Pass	Pass	Fail	Pass
persistent action feedback	3.0	Pass	Pass	Fail	Pass
status text without color	1.0	Pass	Pass	Fail	Pass
unavailable action guidance	1.0	Pass	Fail	Fail	Fail
theme switch preserves workspace	1.0	Pass	Pass	Fail	Pass
comfortable touch targets	1.0	Pass	Pass	Fail	Pass
reduced motion preference	1.0	Pass	Pass	Fail	Pass

Visual

Criterion	Weight	Oracle	GPT	Gemini	Haiku
visual typography	1.0	5/5	5/5	0/5	5/5
visual color and contrast	1.0	5/5	5/5	0/5	2/5

Criterion	Weight	Oracle	GPT	Gemini	Haiku
visual spacing and layout	1.0	5/5	5/5	0/5	5/5
visual hierarchy and scanability	1.0	5/5	5/5	0/5	5/5
visual overall craft	1.0	5/5	4/5	0/5	4/5
visual responsive consistency	1.0	5/5	5/5	0/5	5/5

Recorded non-full-credit explanations

Explanations below are the judges' recorded observations. They are not a claim that all listed subchecks were independently reproduced during packaging. The Gemini shared-gate failure is summarized once to avoid repeating the same root cause for all 86 rows.

GPT-5.4-mini (high)

Functional

draft edit (weight 0.5): The real visible edit request returned 404 and no accepted content/revision change was observed.

ballot open transition (weight 0.5): A clean Verifier room schedule opened, but the required 2-to-3 transition after the intended edit was not obtained.

open definition lock (weight 0.5): No successful edit counterpart existed, so a valid Open edit-lock probe could not be established.

staff write role boundaries (weight 0.5): Create/Open/Close/Publish/membership refusals were observed for non-staff, but no valid successful edit family counterpart existed.

ballot target and revision validation (weight 0.5): The required unknown-target and malformed edit/Open revision tables were not completed with valid edit fixtures.

single choice submission limits (weight 1.0): The required empty and two-choice probes on Leila's open Courtyard were not captured.

accepted votes preserve ballot revision (weight 0.75): The Owen vote exchange was evidence-missing, so the required complete before/vote/after packet set for all three votes was incomplete.

exact vote success receipt (weight 0.75): The complete original Owen Courtyard request/receipt was not captured; later Leila replay evidence cannot substitute for it.

operation id mismatch safety (weight 1.5): Only the Partial approval same-operation mismatch was captured; the required Owen and same-family Create mismatch cases were missing.

approval selection limits (weight 1.5): Malformed approval selection probes were captured on Partial, not on the required untouched Open Verifier ballot.

stale ballot revision refusal (weight 1.0): No accepted edit shape existed from which to perform the required stale-only Draft edit.

published terminal lock (weight 1.0): Published Open/Close/Publish attempts returned 409 unchanged, but the required valid edit counterpart was unavailable.

business audit records (weight 0.75): Create/Open/Close/Publish/membership audit events were observed, but the required accepted edit event was unavailable.

staff mutation success receipts (weight 2.0): Several accepted upstream staff exchanges were retained, but the full seven-action retry/replay ledger was not completed.

refusal receipt after state change (weight 2.0): The dedicated stale-edit and premature-publish refusal rehearsal with later-state replay was not completed.

user wide operation namespace (weight 2.0): Cross-person same-operation votes and a round/Create collision were observed, but the required Create-to-edit and Create-to-membership collision set was incomplete.

durable pending staff work (weight 4.0): Create/Open/Close/Publish/membership reminders survived refresh/restart, but no valid interrupted Edit family was available.

immutable pending retry (weight 4.0): Pending entries survived, but visible Retry produced no request and did not resolve an entry.

independent pending actions (weight 4.0): Q1/Q2/Q3 remained individually visible, but the required independent Retry/Dismiss resolution sequence did not work.

pending actor isolation (weight 4.0): Ruth reminders disappeared for Arun, Leila, and Owen and returned to Ruth, but the required revoked-session recovery refusal was not completed.

cross tab pending resolution (weight 4.0): Same-profile reminders propagated, but dismissal, concurrent ownership, and abandoned-tab retry resolution were not successfully demonstrated.

disjoint draft review (weight 3.0): The required visible merge-review workflow was not available.

conflicting draft review (weight 3.0): The required conflicting-field review and explicit resolution workflow was not available.

draft review rechecks revision (weight 3.0): No valid reviewed-save/recheck workflow was completed.

draft review discard (weight 1.5): No valid refused-edit review discard workflow was completed.

draft review lifecycle stop (weight 1.5): No valid refused-edit lifecycle-stop review workflow was completed.

late recovery reply isolation (weight 2.0): The held late-response account-change sequences were not completed.

round draft conflict (weight 4.0): A selected-ballot conflict refusal was observed, but the required re-review and fresh confirmation sequence was not completed.

round roster conflict (weight 4.0): The three required roster-version race refusals and later fresh confirmation were not completed.

round input boundaries (weight 2.0): Repeated/singleton/unknown, revision-type, and roster-version probes returned 400 unchanged, but the no-active-members and roster-member omission cases were missing.

round refusal receipts (weight 4.0): Only a draft-conflict refusal was available; the required three roster refusals and durable replay set were missing.

round pending recovery (weight 4.0): No accepted interrupted round confirmation was retained as a pending whole-round entry.

Polish

unavailable action guidance (weight 1.0): Draft, Open, and Published cards omit invalid lifecycle actions without explaining why; Closed only explains that results remain hidden until Published. Observer and recorded-member views are read-only, but lifecycle guidance is incomplete.

Visual

visual overall craft (weight 1.0): A coherent design language carries across all reviewed surfaces with rounded outlined cards, accent rails, badges, progress bars, styled empty states, and theme variants. Minor polish issues include repeated tiny empty outline markers on card edges and long repetitive lists that feel slightly under-resolved.

Claude Haiku 4.5

Render

working ballot journey (weight 1.0): The visible Create ballot workflow repeatedly produced POST /api/ballots 400 responses ('A valid request body is required.');

subsequent protected collection reads showed no new ballot. Therefore no genuine five-write journey, vote, close/publish, or persisted result could be completed.

Functional

draft input validation (weight 0.75): Malformed create cases were refused, but the valid UI creation returned 400, so the required positive control was missing.

draft creation (weight 0.75): The valid visible Create request returned 400 and created no ballot.

draft edit (weight 0.5): The valid visible edit request returned 500 with no accepted revision/content change.

ballot open transition (weight 0.5): The valid visible Open request returned 500 and remained Draft.

open definition lock (weight 0.5): No successfully opened Ruth-created ballot existed for a valid lock probe.

fixed eligibility snapshot (weight 1.5): Courtyard seed eligibility was correct, but required Verifier and Future snapshots could not be created.

staff write role boundaries (weight 0.5): Arun and both Members received 403 for all six observed action families, but no accepted Ruth positive counterpart was observed.

ballot target and revision validation (weight 0.5): Unknown targets and several revision types were refused, but array/boolean probes produced 500 and no complete positive counterpart existed.

membership input validation (weight 0.5): Malformed revisions returned 400, but the valid membership mutation returned 500 and no positive control existed.

single choice submission limits (weight 1.0): Empty and two-choice submissions were refused, but the required accepted positive vote failed 500.

private final vote confirmation (weight 1.0): The privacy/final-selection notice appeared, but none of the required accepted votes completed.

accepted votes preserve ballot revision (weight 0.75): No accepted vote packets existed.

exact vote success receipt (weight 0.75): Owen's original Courtyard vote failed 500; no success receipt existed.

approval order independent receipt (weight 0.75): No accepted approval vote or Partial fixture existed.

operation id mismatch safety (weight 1.5): No accepted vote or Create receipt existed to vary and replay.

approval selection limits (weight 1.5): The Verifier approval fixture could not be created or opened.

duplicate participation rejection (weight 1.0): No accepted Courtyard participation existed.

staff identified turnout (weight 0.5): Seed turnout was visible, but required post-vote checkpoints were unavailable.

member participation isolation (weight 0.5): Member reads exposed only own flags, but required accepted-vote stage comparisons were unavailable.

ballot close transition (weight 0.375): Visible Courtyard Close returned 500.

closed vote refusal (weight 0.375): A closed-ballot vote returned 409, but the required clean Future open-then-close fixture was unavailable.

unpublished results hidden (weight 0.75): Current results hid unpublished totals, but the required Courtyard Open and Closed snapshots were incomplete.

publish transition (weight 0.5): Visible Shared-space Publish returned 500.

published approval tally (weight 1.0): Required Shared-space and Partial publication/tally workflows were unavailable.

stale ballot revision refusal (weight 1.0): No accepted editable Verifier draft existed.

business audit records (weight 0.75): Seed audit was faithful, but no accepted edit event was produced in this run.

staff mutation success receipts (weight 2.0): Observed staff mutations returned 400/500, so the required seven success receipts and retries were unavailable.

roster conflict snapshot chain (weight 2.0): Membership writes returned 500; no two-tab race or snapshot chain was possible.

refusal receipt after state change (weight 2.0): No valid authorized stale or premature refusal receipt was established.

user wide operation namespace (weight 2.0): No accepted Create or vote receipts existed for namespace checks.

durable pending staff work (weight 4.0): Failed actions left reminders, but not intentionally interrupted accepted upstream responses.

immutable pending retry (weight 4.0): No accepted upstream receipt or complete Retry identity comparison existed.

independent pending actions (weight 4.0): Failed-action reminders did not establish the required accepted Q1/Q2/Q3 sequence.

pending actor isolation (weight 4.0): No qualified accepted Ruth pending work was established.

cross tab pending resolution (weight 4.0): No qualified shared-profile accepted queue sequence was completed.

disjoint draft review (weight 3.0): No valid second Draft or review-save baseline existed.

conflicting draft review (weight 3.0): No valid shared Draft conflict workflow existed.

draft review rechecks revision (weight 3.0): No refused valid Draft edit reached review confirmation.

draft review discard (weight 1.5): No review workspace was reachable.

draft review lifecycle stop (weight 1.5): No successful editable/opened Draft review existed.

late recovery reply isolation (weight 2.0): No accepted interrupted Ruth actions existed for account-change races.

round review and cancel (weight 3.0): No visible round workflow or valid Draft pair existed.

round atomic open (weight 4.0): No round confirmation workflow existed.

round draft conflict (weight 4.0): No round review/confirmation control existed.

round roster conflict (weight 4.0): Membership writes failed and no round workflow existed.

round input boundaries (weight 2.0): No genuine round confirmation request was observed.

round role boundaries (weight 1.0): No genuine round confirmation shape or accepted positive existed.

round success receipt (weight 4.0): No accepted round receipt existed.

round refusal receipts (weight 4.0): No round refusal receipts existed.

round pending recovery (weight 4.0): No accepted round confirmation could be interrupted and retained.

Polish

semantic labels and landmarks (weight 1.0): Sign-in and voting controls are labeled, with main/navigation landmarks and headings. The New ballot form has visible generic labels but no associated accessible names for Title, Description, Voting Method, or Choices.

unavailable action guidance (weight 1.0): Lifecycle cards expose only next actions without disabled states or explanations for omitted invalid actions. Arun's Observer Members view exposes Pause/Activate, and Leila's recorded view only says "You participated" without explicit submission guidance.

Visual

visual color and contrast (weight 1.0): Light theme is cohesive and readable, but the dark theme has substantial contrast problems: green counts and outline-action text nearly disappear, and the pale Results warning has almost invisible white text.

visual overall craft (weight 1.0): A coherent component language carries through cards, tables, status chips, buttons, results, members, and the form. Minor polish issues remain in sparse secondary views and dark-theme secondary controls.

Gemini shared-gate failure

The shared gate failed: the visible correct-login attempt received HTTP 400 with “Unexpected field: operation_id.” No authenticated Ruth workspace was established, so all criteria are invalidated.

The complete per-criterion reasons, including all Gemini gate-zero verdicts and all Oracle full-credit reasons, are preserved in each job archive under verifier/reward-details.json and verifier/judges/<dimension>/attempt-0001/final.json.

Delivery and evidence index

The task ZIP is byte-identical to frozen r27, with 29 files matching current source. All five exported Oracle application files match the golden solution. Every trial has the same recorded task checksum and Harbor digest. All 12 exported verifier hashes (five prompts, five judges, runner and reward configuration) match the delivered task.

Task ZIP SHA-256: 76b8cceb5eb22b2a48df6d53d58f647814516722b4beed6e8f7d7fef625cc578

Recorded task checksum: 9c16b198807d890250b03d4c1c56efb8870ca6bd0c4de9b4d6e59c5cdf5b8210

Recorded Harbor digest: sha256:a1af9ceddaaf204c997d72683a69cca151b7d6f632efb7f43da533be054639a8

These are different identifier schemes; the platform checksum and digest are reported from the exports. The four named job archives preserve all supplied run files; Oracle includes its NOP sibling. Detected provider credentials, if present, are redacted only in delivery copies and recorded in the external package audit. Scores, verdicts and original exports are preserved.

Oracle: run-07a88cf3-85a8-49a9-ba27-58e30a367e71 / common-ground-ballot__iBeMVH2

GPT-5.4-mini (high): run-5980ca27-03f9-4f79-976a-dc1349d3c8c8 / common-ground-ballot__3pecCc8

Gemini 3.7 Flash: run-7079a467-6ff1-4235-9f93-45b3c35b356b / common-ground-ballot__97FPuUp

Claude Haiku 4.5: run-1b99db31-d70f-4b81-9290-acccc11908d8 / common-ground-ballot__pWxzQKF

NOP: run-07a88cf3-85a8-49a9-ba27-58e30a367e71 / common-ground-ballot__cQy8rSn, included in the Oracle job ZIP.

Each job ZIP includes the original run-level metadata, trial results, agent trajectories, submitted application and verifier evidence present in its supplied export. No submission is repaired or replaced. The companion case study explains the product and the observed separation; this evaluation report provides the score ledger and limitations.

Pre-delivery local record: 101 golden browser checks, eight detected round mutants, 30 runner checks and 294 ZIP checks. Packaging rechecks ZIP integrity, exact source/provenance hashes, score arithmetic, complete verdicts and document readability. It does not rerun those historical browser tests, rebuild Docker images, launch paid evaluations or perform platform QC.