

Evaluation Report — Gambit Hollow Cribbage

Oracle, GPT, Gemini and interrupted Haiku evidence; no score changes

turing/gambit-hollow-cribbage | version 1.0.0 | 15 September 2026

Result

Oracle and GPT satisfy the saved overall requirements. Oracle is 0.9833 with all 40 Functional criteria passing; GPT is 0.6616. Gemini is above 0.7, Haiku needs a clean rerun, and Oracle Visual remains below the preferred 1.0. This is the final package of the supplied evidence, not a claim that every delivery prerequisite is complete.

Run	Overall	Functional	Polish	Visual	Functional count
Oracle	0.9833	1.0000	1.0000	0.9167	40/40
GPT-5.4-mini (high)	0.6616	0.7083	0.6000	0.5833	29/40
Gemini 3.7 Flash (high)	0.8134	0.9167	0.4000	0.9167	37/40
Claude Haiku 4.5	0.0000	0.0000	0.0000	0.0000	0/40; interrupted
NOP	0.0000	0.0000	0.0000	0.0000	No-op control

Render and Constraints are 1.0 for Oracle, GPT and Gemini. Haiku has both gates zero. NOP is graded=0, no_op=1; its zeros are control outputs.

Task and verification contract

Area	Coverage
Exact scoring	Forty supplied hand fixtures; overlapping fifteens, pairs and runs; flush/nobs rules; a live hand and pegging scoring bench with invalid-input rejection.
Complete cribbage match	Two seats, six-card deals, two discards each, hidden hands/crib, cut, legal pegging, go/31 resets, ordered show, alternating dealers and instant capped wins at 121.
Reproducible practice	Named deals and near-target starts exercise specific scoring boundaries, a seven-hand club match and continuation when one player has exhausted their cards.
Durable recovery	SQLite saved games and club ladder; revision checks, accepted-request identity binding, exact retry replay, lost-response UI recovery and two real process restarts.
Usable presentation	Inline SVG board/cards, real browser controls, keyboard focus, mobile reachability, pending-action feedback and six Visual assessment axes.

Node/Express application at localhost:3000, started with node serve.js; SQLite at /app/gambit.db. Inline SVG board and cards. The verifier uses Codex, gpt-5.6-luna, max reasoning and temperature 0, with pinned Playwright tools. Agent budget: 7200 seconds; separate verifier budget: 13200 seconds. Public agent/verifier networking is retained.

Dimension	Criteria	Budget	Scoring
Render	2	600 s	all_pass gate
Constraints	2	600 s	all_pass gate
Functional	40	9000 s	weighted mean: 60%
Polish	5	900 s	weighted mean: 20%
Visual	6	900 s	weighted mean: 20%

If Render or Constraints is zero, final reward is zero. Otherwise final reward = 0.6 × Functional + 0.2 × Polish + 0.2 × Visual, rounded to four decimals using the exported dimension scores. Functional counts are unweighted; Visual raw 4/5 maps to normalized 0.75 and 5/5 maps to 1.0.

Verifier elapsed time

Run	Verifier elapsed	Trial status
Oracle	33m 43s	No recorded trial exception
GPT-5.4-mini (high)	41m 41s	No recorded trial exception
Gemini 3.7 Flash (high)	35m 05s	No recorded trial exception
Claude Haiku 4.5	5m 27s	Agent exit 143; partial artifact graded

Findings

Oracle

Overall 0.9833, Functional 1.0000 with all 40 criteria passing, and Render/Constraints/Polish 1.0000. Visual 0.9167 reflects two raw 4/5 verdicts: spacing/layout and responsive consistency. Reasons cite blank spacing transitions, dense mobile controls and a shortened practice label. The required overall threshold passes, while the preferred Visual=1.0 authoring target is not reached.

GPT-5.4-mini

Overall 0.6616 is within the saved 0.1–0.7 range. Functional is 0.7083: 29/40 criteria passed, earning 34/48 weighted points. Eleven failures concern pegging/run/last-card calculations, a repeated discard accepted out of order, show totals, hand-state reset, match completion and lost-response retry after reload. Polish loses keyboard focus after a card selection and overflows at desktop/mobile widths. The basic working-table gate passed, making this more informative than the earlier two browser-blocked zero attempts.

Gemini 3.7 Flash

Overall 0.8134 is above 0.7. Functional is 0.9167, with 37/40 criteria passed and 44/48 weighted points. Failures concern invalid starting scores accepted/clamped, a duplicated physical card accepted by the pegging bench and loss of pending retry after reload. Polish loses mobile reachability, an explanatory illegal-move refusal and duplicate-action suppression. No separate Gemini acceptance band was found; this run would fail an all-model ceiling of 0.7.

Claude Haiku 4.5

The exported reward is 0.0000 with graded=1 and no_op=0, but the agent did not finish normally. UnknownApiError records command exit 143 immediately after `kill -f "node serve.js"`. The broad pattern likely also matched the agent command line, which contains the task instruction; self-termination is an inference, not a demonstrated process trace. The partial app returned "UI coming soon" from create/open controls. Every dimension assigned zero after the shared playable-table gate failed. Preserve this grade and rerun Haiku for clean completed-model evidence.

Limitations and remaining evidence

- This delivery packages existing supplied exports. It does not change task files, adjust recorded verdicts, rerun models or certify new platform QC.
- The saved GPT band applies to overall reward. Its Functional 0.7083 is slightly above 0.7, and the earlier aspirational overall target near 0.4 was not reached.
- The two Oracle Visual deductions remain in the report. Overall acceptance does not imply a perfect visual result.
- Haiku is interrupted evidence, not a normally completed attempt. Its 55 gate-zero verdicts are not 55 independently exercised implementation defects.
- Shared setup or gameplay failures can prevent later checkpoints. Exported failed criteria must not all be treated as independent root causes. This review did not replay every model browser action.

- Only one current attempt per model is included. These outcomes establish neither stable rankings nor repeatability across new runs.
- The supplied model set is GPT, Gemini and Haiku. No Sonnet export is available; Gemini is not relabeled as Sonnet. If the general checklist requiring Sonnet still applies, that evidence remains outstanding.
- No platform static/rubric QC report was supplied. Local archive/configuration checks are reported separately. The earlier fresh local environment build was blocked by the package proxy; no new local Docker build is claimed.

All 55 criterion outcomes

Binary Pass/Fail and raw Visual /5 values reproduce the supplied verdicts. Haiku failed the common browser prerequisite; its rows do not establish that every individual behavior was exercised.

Render

Criterion	Weight	Oracle	GPT	Gemini	Haiku
board renders	2.0	Pass	Pass	Pass	Fail
cards render as cards	1.0	Pass	Pass	Pass	Fail

Constraints

Criterion	Weight	Oracle	GPT	Gemini	Haiku
server backed local application	2.0	Pass	Pass	Pass	Fail
drawn as svg not images	1.0	Pass	Pass	Pass	Fail

Functional

Criterion	Weight	Oracle	GPT	Gemini	Haiku
seed ladder import	1.0	Pass	Pass	Pass	Fail
the twenty nine hand	1.0	Pass	Pass	Pass	Fail
the app agrees with the supplied fixture	1.0	Pass	Pass	Pass	Fail
a card counts in every combination	1.0	Pass	Pass	Pass	Fail
the crib flush needs five	1.0	Pass	Pass	Pass	Fail
a five card flush counts in both	1.0	Pass	Pass	Pass	Fail
the ace is low and does not wrap	1.0	Pass	Pass	Pass	Fail
nobs needs the matching suit	1.0	Pass	Pass	Pass	Fail
fifteen in the play is two	1.0	Pass	Fail	Pass	Fail
a run in the play need not be in order	1.0	Pass	Pass	Pass	Fail
the count cannot pass thirty one	1.0	Pass	Pass	Pass	Fail
a hand that is not a hand is refused	1.0	Pass	Pass	Pass	Fail
two cards each go to the crib	1.0	Pass	Pass	Pass	Fail
the crib is hidden until the show	1.0	Pass	Pass	Pass	Fail
moves out of order are refused	1.0	Pass	Fail	Pass	Fail
the show counts in order	1.0	Pass	Fail	Pass	Fail
the game stops at the target	2.0	Pass	Pass	Pass	Fail
a game is a sequence of hands	2.0	Pass	Fail	Pass	Fail
the deal alternates between hands	1.0	Pass	Pass	Pass	Fail
a game survives a reload	1.0	Pass	Pass	Pass	Fail
the ladder updates when a game ends	1.0	Pass	Pass	Pass	Fail
hand order invariance	1.0	Pass	Pass	Pass	Fail
pairs and last card in real play	1.0	Pass	Fail	Pass	Fail
go resets and other player leads	1.0	Pass	Pass	Pass	Fail
heels immediate and capped	1.0	Pass	Pass	Pass	Fail
show wins stop at each boundary	1.0	Pass	Fail	Pass	Fail
practice and scoring do not corrupt saved games	1.0	Pass	Pass	Fail	Fail
pegging scoring bench	1.0	Pass	Fail	Fail	Fail
runtime manifest routes	1.0	Pass	Pass	Pass	Fail
final restart persistence	1.0	Pass	Pass	Pass	Fail
complete club match	2.0	Pass	Fail	Pass	Fail
one player exhausted pegging	1.0	Pass	Fail	Pass	Fail
last card thirty one	1.0	Pass	Fail	Pass	Fail
independent saved games	1.0	Pass	Pass	Pass	Fail

Criterion	Weight	Oracle	GPT	Gemini	Haiku
stale tabs and revisions	2.0	Pass	Pass	Pass	Fail
durable creation identity	2.0	Pass	Pass	Pass	Fail
accepted action replay	2.0	Pass	Pass	Pass	Fail
accepted identifier binding	1.0	Pass	Pass	Pass	Fail
lost response retry ui	2.0	Pass	Fail	Fail	Fail
recovery phase matrix	2.0	Pass	Pass	Pass	Fail

Polish

Criterion	Weight	Oracle	GPT	Gemini	Haiku
keyboard controls and focus	1.0	Pass	Fail	Pass	Fail
responsive controls and reachability	1.0	Pass	Fail	Fail	Fail
illegal move feedback	1.0	Pass	Pass	Fail	Fail
pending action and empty state	1.0	Pass	Pass	Fail	Fail
readable state and reduced motion	1.0	Pass	Pass	Pass	Fail

Visual

Criterion	Weight	Oracle	GPT	Gemini	Haiku
visual typography	1.0	5/5	4/5	5/5	0/5
visual color and contrast	1.0	5/5	3/5	5/5	0/5
visual spacing and layout	1.0	4/5	3/5	5/5	0/5
visual hierarchy and scanability	1.0	5/5	3/5	5/5	0/5
visual overall craft	1.0	5/5	4/5	5/5	0/5
visual responsive consistency	1.0	4/5	3/5	3/5	0/5

Recorded non-full-credit reasons

These are exported judge explanations. They are observations from this run, not independently reproduced causal diagnoses.

Oracle

Visual: visual spacing and layout

Desktop board/sidebar grouping, card rows, score tiles, show dividers, and mobile stacking are well organized; a few larger blank transitions and compressed mobile controls are minor imperfections.

Visual: visual responsive consistency

At 375px the palette, typography, cards, show, ladder, and bench remain coherent and stack naturally without visible horizontal breakage; dense controls and a shortened practice-deal label show minor mobile compression.

GPT-5.4-mini (high)

Functional: fifteen in the play is two

The 15 case awarded two, but the required 29-to-31 A play awarded five instead of exactly two.

Functional: moves out of order are refused

A fresh second discard by A was accepted before B discarded.

Functional: the show counts in order

Entries were ordered B/A/crib, but B scored 4 instead of the required 6 and final scores diverged.

Functional: a game is a sequence of hands

Next hand advanced identity and dealer state, but prior played cards remained and show totals were wrong.

Functional: pairs and last card in real play

Early pair checkpoints matched, but the final card omitted the required pair and last-card points.

Functional: show wins stop at each boundary

The three boundary games diverged before reaching the required show winner boundaries.

Functional: pegging scoring bench

The required five-card pile plus 2S returned five points instead of two.

Functional: complete club match

The first hand totals diverged and a later expected legal card could not be played, so the seven-hand match did not complete.

Functional: one player exhausted pegging

The exhausted-player journey produced incorrect run/go scoring and ended with B pegging 8 instead of 1.

Functional: last card thirty one

Initial 31 and run checkpoints worked, but final pegging/show totals diverged.

Functional: lost response retry ui

Intercepted discard was accepted server-side and Retry appeared, but reload lost the pending receipt so post-reload retry was unavailable.

Polish: keyboard controls and focus

Controls had accessible names and visible Tab/Shift+Tab focus; Enter switched seats and Space selected a card. After selecting 5♥ with Space, focus reset to BODY instead of a useful control.

Polish: responsive controls and reachability

Horizontal overflow was present at both sizes: scrollWidth 1359 at 1280px and 606 at 375px. The fixed-width ladder/history tables extend off-screen.

Visual: visual typography

Clear type scale distinguishes the title, section labels, seats, totals, controls, bench, and show breakdowns. Minor issues include small secondary labels and wrapped/clipped ladder/history identifiers.

Visual: visual color and contrast

The forest-green, mint, yellow, cream-card palette is cohesive, with distinct red/black suits and peg tracks. Open-game IDs and some secondary history text are visibly low-contrast on green panels.

Visual: visual spacing and layout

Board metrics, peg rail, seat cards, card grid, show panels, bench, and ladder use deliberate grouping. The long open-game list creates excessive vertical separation, while ladder/history tables overflow and clip content.

Visual: visual hierarchy and scanability

The active board has a clear metric row, seat grouping, turn panel, show breakdown, and commentary. However, the initial desktop view is dominated by many open-game rows, pushing the primary game state far below.

Visual: visual overall craft

A consistent board-game component language carries through SVG cards, peg board, panels, controls, empty states, scoring bench, and show history. Overflowing ladder content, raw history JSON, and low-contrast IDs prevent a top score.

Visual: visual responsive consistency

At 375px, the title, board metrics, cards, show breakdown, and bench stack coherently while retaining the desktop palette and styling. The ladder retains desktop table width, causing horizontal overflow and clipped columns/history summaries at the narrow viewport.

Gemini 3.7 Flash (high)

Functional: practice and scoring do not corrupt saved games

Invalid starting scores were incorrectly accepted by direct create requests, clamped, and created additional games instead of being refused.

Functional: pegging scoring bench

All listed experiments worked except duplicate [5C]/5C, which incorrectly returned 200 with pair points instead of refusing.

Functional: lost response retry ui

Forward-and-abort proved the server accepted the discard and the UI initially showed Retry Move, but after the required reload the retry control disappeared, so safe UI recovery could not be completed.

Polish: responsive controls and reachability

At 375×760 the document overflowed horizontally (scrollWidth 424 vs 375); saved-game tables extended to about 803px and ladder history was offscreen. Desktop 1280×800 fit.

Polish: illegal move feedback

At count 28, 9D was disabled as “unplayable,” but had no textual reason, title, description, or live announcement explaining that it would exceed 31. Valid AC still worked.

Polish: pending action and empty state

The no-game/practice view had a usable “Start Game at Table” control. However, while the first game POST remained pending, repeated mouse and keyboard activation produced another accepted 201 response and duplicate unfinished game IDs.

Visual: visual responsive consistency

The mobile table, navigation, cards, and scoring bench adapt coherently, but the ladder remains a wide two-column desktop layout with match history beginning offscreen, and the mobile show panel slightly exceeds the viewport. This creates visible horizontal overflow and loss of secondary-surface hierarchy at 375px.

Claude Haiku 4.5

All 55 verdicts are zero following the shared playable-table prerequisite. Representative Render explanation: The local page and ladder request loaded, but the shared gate failed: the visible Board is only a static score header/control shell, and New Game/Practice Deal displayed “coming soon” alerts. No substantive peg board or working game could be created.

The complete individual explanations remain in this run archive at verifier/reward-details.json. Repeating all 55 prerequisite failures here would not add evidence of independent defects.

Evidence index and package identity

Oracle: run-77e08fa0-1b6c-4e21-ad8c-6b485409f7bd / gambit-hollow-cribbage__TuEG7Z5

GPT-5.4-mini (high): run-e4be6735-badc-4ed1-83e5-0749b43d0074 / gambit-hollow-cribbage__cmZxA5f

Gemini 3.7 Flash (high): run-8b309ea5-846a-4cb9-82d7-5d4b95e877bb / gambit-hollow-cribbage__q9jV28

Claude Haiku 4.5: run-5eaceacd-12e2-4035-9105-facdeeeaa271 / gambit-hollow-cribbage__WGZSAcM

NOP: run-77e08fa0-1b6c-4e21-ad8c-6b485409f7bd / gambit-hollow-cribbage__RKrf7iu

The task ZIP is byte-identical to the frozen browser-acceptance candidate. All 39 current task files match it; every golden solution file matches the Oracle app export. Every graded run contains the same 12 verifier hashes (five prompts, five judge files, runner and reward configuration). All five trials share the recorded task checksum and Harbor digest below.

Task ZIP SHA-256: 06a4f99349469c181d51b8da5520fa0c8bc02332c0433cbfc0d826b374e57f6a

Recorded task checksum: 770ea7287ee52452d6f44b7523b120e216b1a4d080969516efdeecdd8c70837c

Recorded Harbor task digest: sha256:ebe8a7bc5b75c6b6a8151c8cb985d85070042698e2888bc48ded8f6a7c230bb8

These three identifiers are different hash schemes and are not interchangeable. The platform checksums are reported from exports, not recomputed from the task ZIP. Original run exports are unchanged; archive copies preserve scores and verdicts and redact detected provider credentials if present. The Oracle archive includes its original NOP sibling.

Final packaging repeats local archive/configuration checks and compares frozen task bytes, provenance, criterion counts, reward arithmetic and original artifact hashes. It does not rerun browser suites, Docker builds, paid models or platform QC. Detailed local packaging checks are in the accompanying submission-preparation directory.